Pre-Flight

Tokens burned. Time lost. Nothing logged.

Sonakshi Pandey

THE PROBLEM

Tiffin runs 6 AI agents calling MCP tool servers through the Databricks AI Gateway.

When a tool rejects a semantically reasonable argument the agent does not fail loudly - it retries, reshapes the argument, guesses, and burns tokens.

These failures return HTTP 200, so conventional error dashboards never see them.

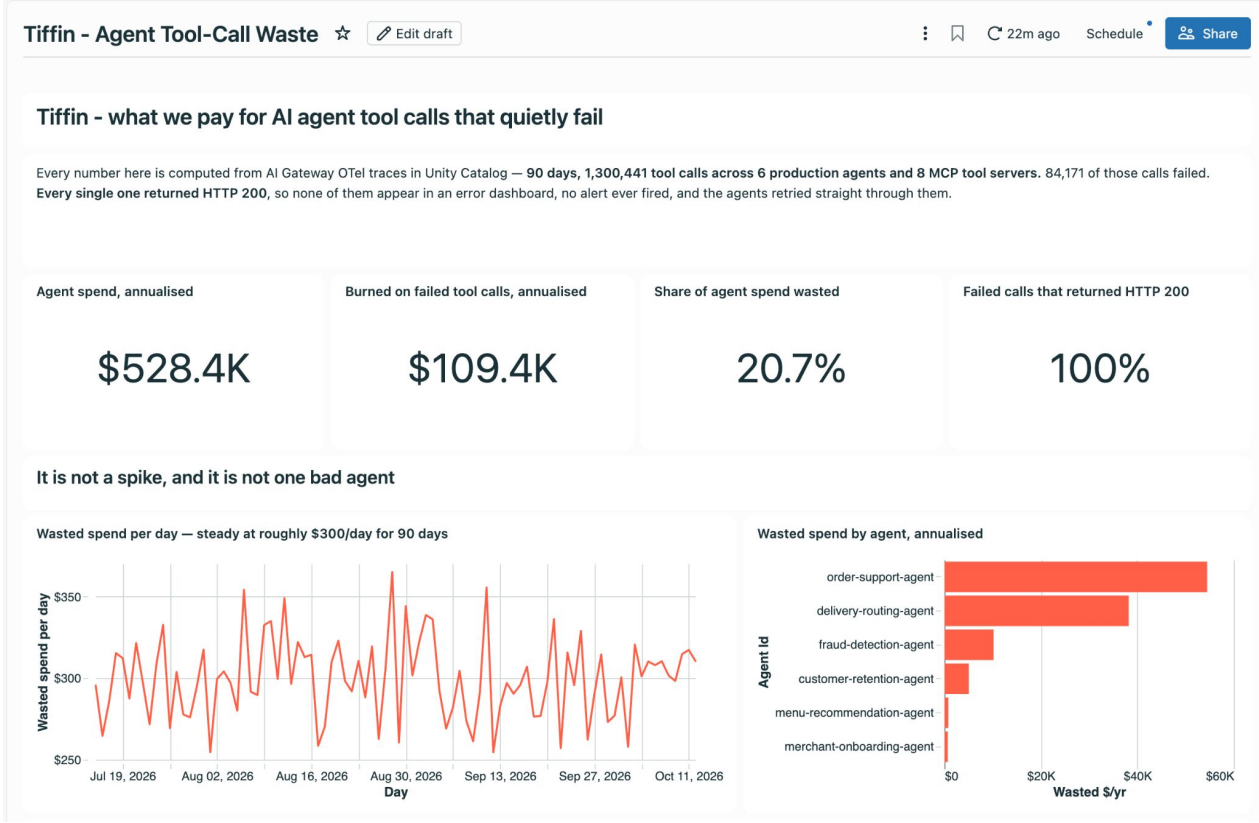
THE SOLUTION - Stop the tokens and time our agents burn on tool calls that fail or retry

A guard model scores every tool call before dispatch and returns a failure probability plus a verdict: coerce or allow. On coerce the agent rewrites the bad argument from the server's own declared schema, so the failure and its retry loop are prevented at dispatch, not diagnosed afterwards.

VP Engineering: 20.7% of agent spend recovered, agents right instead of confidently wrong. Tool-server owner: knows which contract to repair first, ranked by dollars rather than error count, because cost per failed call varies 102x.

TIFFIN

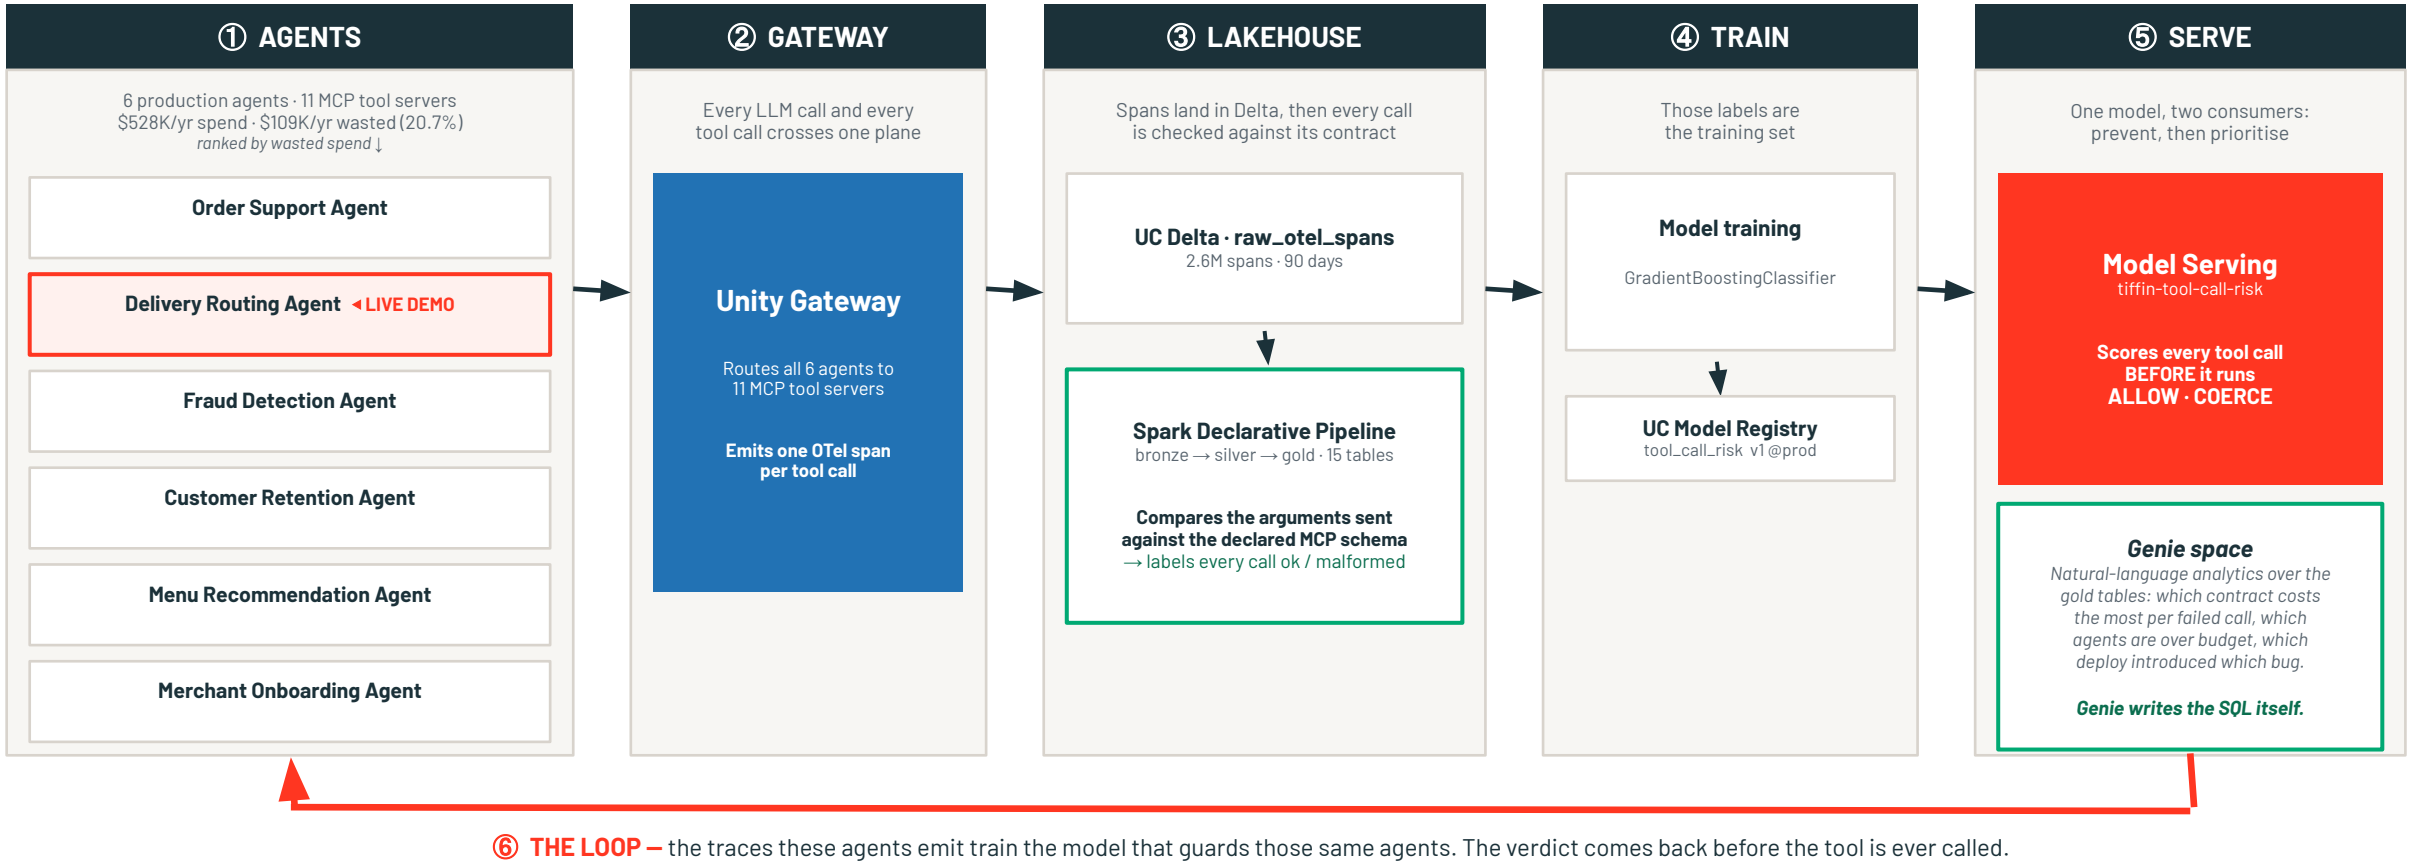
Food-delivery operations · 6 AI agents · 11 MCP tool servers · 109K wasted spend



©2026 Databricks, Inc. — All rights reserved

Six agents. One Gateway. One closed loop.

All six production agents route through the AI Gateway → it emits an OTel span for every tool call → Lakeflow compares the arguments sent against the declared MCP schema → those labels train the model → the model scores the next call before it runs.



⑥ **THE LOOP —** the traces these agents emit train the model that guards those same agents. The verdict comes back before the tool is ever called.

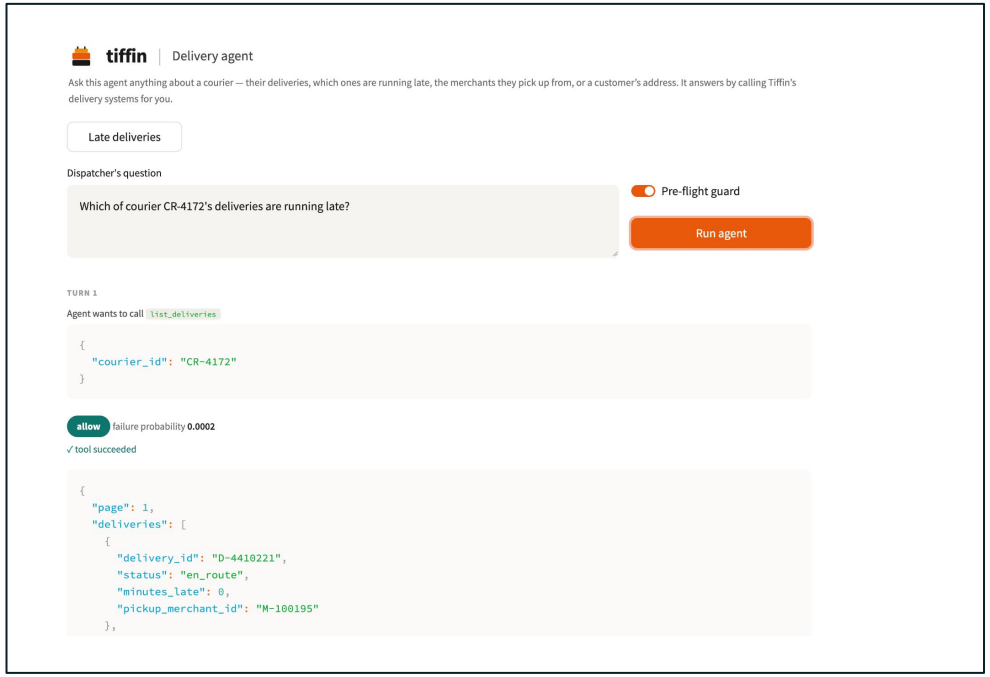
Unity Catalog — one governance plane across all of it: raw spans · 15 pipeline tables · model lineage · serving endpoint · Genie space

©2026 Databricks, Inc. — All rights reserved

Links



Genie Space



Delivery Agent

REAL — functional today on e2-dogfood
Gateway OTel schema · Lakeflow pipeline (15 tables) · MLflow model + leave-one-bug-family-out eval · UC registry · Model Serving endpoint · Genie space · Databricks App · AI/BI dashboard
github.com/29sonakshi/tiffin-preflight — code + committed execution evidence

MOCKED — only to make the story demoable
Synthetic Tiffin trace corpus (2.6M spans generated from a seeded bug registry) · 5 in-process MCP tool mocks inside the demo app

©2026 Databricks, Inc. — All rights reserved